

CSE 2221 SYSTEMS PROGRAMMING

Introduction to system programming.

- Systems programming / system programming is the activity of programming computer system software.

- The primary distinguishing characteristic of system programming when compared to app programming is that application programming aims to produce SW which provides services to the user directly e.g. wordprocessors, information management systems etc. whereas systems programming aims to produce software and software platforms which provide services to other software e.g. operating systems, computational science applications, game engines, videogames, industrial automation etc.

- System programming requires a great degree of hardware awareness its goal is to achieve efficient use of available resources either because the software itself is performance critical e.g. videogames or because even small efficiency improvement directly transforms into significant monetary savings for the service provider.

- The following attributes characterize system:

1. Programmer can make assumptions about HW and other properties of the system that the program runs on and will often exploit those properties for example by using an algorithm that is known to be efficient when used with specific HW.
2. Usually a low level programming language is used so that
 - a) Programs can operate in resource constrained environment.
 - b) Programs written to be efficient with little runtime overhead
 - c) Programs may use direct control over memory access and control flow.
 - d) Programmer may write parts of the program directly in assembly language.
3. Often system programmer cannot be run in a debugger.

In system programming often limited programming facilities are available the use of automatic garbage collection is not common and debugging is sometime hard to do.

N/B: Garbage collection - process of finding and deleting objects which are no longer being referenced by other objects in other words garbage collection is the process of removing any objects which are not being used by any other objects.

Implementing certain parts in operating systems and networking requires system programming example - Implementing paging (virtual memory) or a device driver for an OS

A Programming Illusion.

- System programs can sometimes be written to extend the functionality of the operating system itself and provide functions that higher level application can use.
- The following program use the API of the UNIX as
- The following program gets its input from the keyboard or a disk file and writes its output to display screen or to file on the disk. Such programs are called console applications because the keyboard and the control screen are part of the control devices.
- UNIX automatically creates the standard input and standard output devices for it, opens them and makes them ready for reading and writing respectively.
- In C (and C++), `stdin` and `stdout` are variables in the preprocessor header file that relates to the standard input and standard output devices respectively.
- By default, the keyboard and the display of the associate terminal are the standard input and output devices respectively.

1. Examples.

C program using a simple I/O model.

```
#include <stdio.h>
```

```
/* copy from stdin to stdout */
```

```
int main()
```

```
{
```

```
    int c;
```

→ end of file

```
    while ((c = getchar()) != EOF)
```

```
        putchar(c);
```

```
    return 0;
```

```
}
```

2. Simple C++ program using I/O model.

```
#include <iostream>
```

```
using namespace std;
```

```
/* copy from stdin to stdout using C++ */
```

```
int main()
```

```
{
```

```
    char c;
```

```
    while ((c = cin.get()) & !cin.eof())
```

```
        cout.put(c);
```

```
    return 0;
```

```
}
```

$\text{cin} \gg \rightarrow \text{cin.get}$.

$\text{cout} \gg \rightarrow \text{cout.put}$.

These programs give us the illusion that they are directly connected to the keyboard & the display device via C library functions; `getchar()` and `putchar()`; and the C++ `iostream` member functions `get()` and `put()`. Either of them can be run on a single user desktop computer or on a multiuser time shared workstation in a terminal window and the results will be same.

Process:

A program is an executable file, and a process is an instance of a running program.

- When a program is run on a computer, it is given variable resources such as a primary memory space, mappings of various kinds and privileges such as the right to read or write certain files or devices.

- As a result at any instance of time, associated to a process is the collection of all resources allocated to the running program as well as any other properties and settings that characterize that process, such as the value of the processor's registers.

UNIX assigns to each process a unique number called its process-ID or PID for example; at a given instant of time, several people might all be running the `gcc` (compiler, `gcc`).

- Each separate execution instance of `gcc` is a process with its own unique pid.

- This `ps` command can be used to display which process with its own unique pid.

The person are running and various options to it control what it outputs.

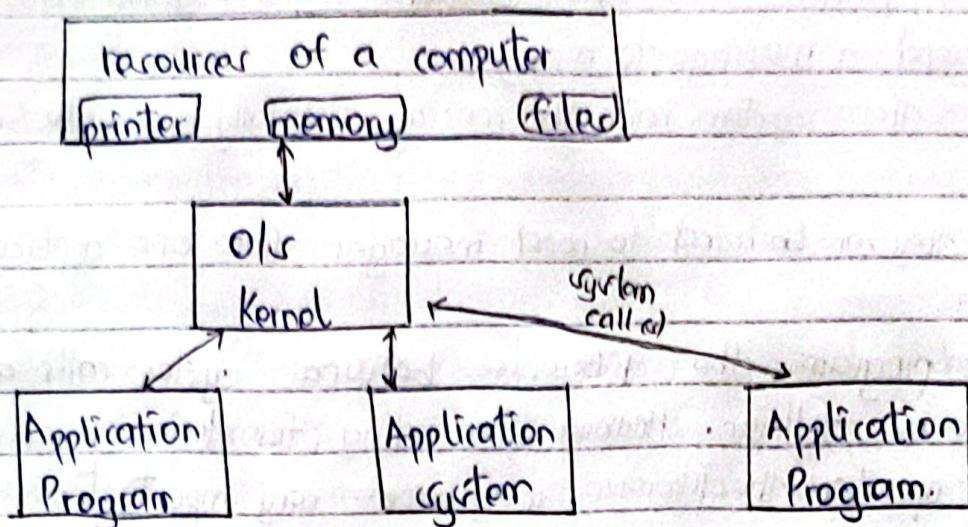
- At the programming level, the function `getpid()` returns the process id of the process that invokes it.

- The following program shows how it is used; we will also see that `getpid()` is an example of a system call

System call

```
#include <sys/types.h>
#include <sys/conf.h>
int main()
{
    printf("I am the process with process-id %d\n", getpid());
    return 0;
}
```

System calls.



An OS is a program with two jobs: it manages the resources of the computer (files, memory, peripherals etc) and it manages all other programs running on the computer.

The programs we write access the resources of the computer by calling functions executed by the kernel.

These functions are called system calls.

A collection of system calls is sometimes referred to as API.

There are dozens of system calls supported by any Linux or Unix kernel.

Some common examples include (none of these lists is all inclusive)

① File I/O : `open()`, `read()`, `write()`, `close()`, `creat()`, `link()`, `lseek()`.

These system calls ask the OS to access a file

② access control : `fork()`, `exec()`, `execv()`, `wait()`, `system()`

These system calls ask the OS to run another program or control how it runs.

③ Memory management : `mmap()`, `brk()` - These system calls ask the OS to provide memory to be used by the application program.

④ Time management: `time()`, `setitimer()`, `gettimeofday()`, `alarm()`. These system calls ask the OS to access the system clock in some cases taking an action that affects a program when a given time interval has passed.

Interprocess communication (IPC).

signal(), pause(), kill(), sigaction() . This system calls asks the OS to send a message to a program.
In most of these system calls the message usually originate in the kernel.

Although they can be used to send messages from one application program to another.

for most programs the differences between system calls and library functions are as follows: However when coding for device drivers or other low level operation the differences may become very important:

Standard library functions are built on top of system calls.
example malloc() and related functions are built on top of the mmap() and brk() system calls, meaning they call them to get the job done.

The malloc function is in the C standard library while mmap is a system call.

This means that greater control over the computer resources can be obtained through system calls.

Standard library functions tend to be more standardized than system calls. They have both standards bodies that is; ANSI defines library standard functions while the POSIX defines system call standard. however since the system calls provide direct access to the kernel different OS's (including different variations of UNIX) will have slightly at least slightly different system calls.

System calls cause a "context switch"

meaning that the application program stops while the kernel program runs for a while to perform the requested operation.

Because of this a system function call takes longer than a normal function call within your own program.
Standard library calls often optimize operations to minimize the number of system calls thus speeding up program execution.

Q) What is a system call; why would a program make one.

Ans) What are the differences between a system call and a library function call.

Ans) Is malloc() a system call or a library call.

Ans) What are the differences between open(), read(), write(), close() and fopen(), fread() and fclose().

Ans) In computing a system call is how a program requests a service from an OS kernel; thus may include hardware related services eg accessing the hard disk(), creating an executing new processes and communicating with integral kernel services like scheduling.

System calls provide an essential interface between a process and the operating system.

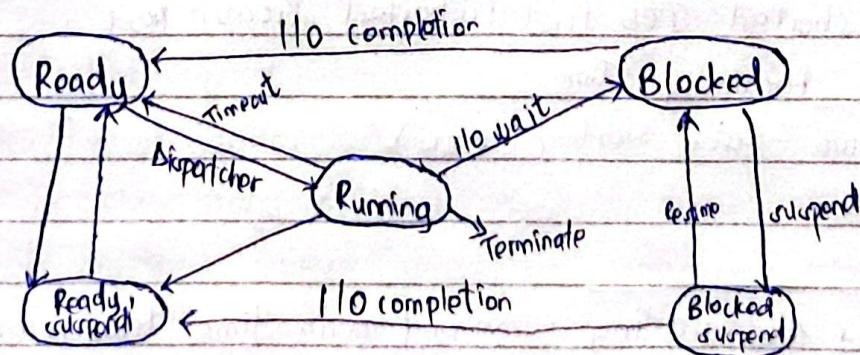
Q)

PROCESS MANAGEMENT AND SCHEDULING.

The major responsibilities of a multitasking OS is process management. The OS must allocate resources to processes enable processes to share and exchange information protect the resources of each process from other processes and enable synchronization amongst processes.

A process can be defined as a program in execution it can be defined as a program currently making use of the processor at any one time.

The diagram below shows the various states of a process:



PCB - process control Block. - u

- Has process state counter

INTER PROCESS COMMUNICATION (IPC)

Interfacing that allow programmers to communicate and allow programs to run concurrently.
Achieved using the following intercommunication

Data to flow from one process to another only in
only from the output of one process to the input of

Communication made possible using 2 pipes; 1 in
which must have a common process origin -
output process to the input process is buffered.
process receiver it.

Reading data from a program and passing it onto

Document from the computer system where re
turn.

is a specific name.

processes that do not have a shared

.

in as First in First Out (FIFO).

Processor to pass messages between
single or several message queues
for each message queue by giving
go on the queue.

created by one process and then used
and or write messages to the queue.
(processor) create message queue

- CPU register such as index register.
- Condition codes.
- accumulators.
- Accounting information such as amount of CPU and real time used.
- Memory management info eg Base and index register
- Input/Output status information

Process Scheduling.

.. long term scheduler

2. medium-term scheduler

3. short-term scheduler

Scheduling Policies

1. Non pre-emptive

2. Pre-emptive

Pre-emptive scheduling permits a process to be removed from the processor when other high priority processes come in. but non preemptive scheduling cannot permit this it works on a first come first serve its fairer though it keeps short jobs waiting.

Scheduling algorithm.

1. FIFO

2. RR- Round Robin.

3. SJF (SPN - shortest Job first /shortest Process Next)

4. SRT- shortest Remaining Time.

5. Higher Response Ratio Next (HRRN)

HRRN what is the waiting time and service time briefly explain and give an example.

What are the scheduling objectives

INTER - PROCESS COMMUNICATION (IPC)

It is a set of interfaces that allow programmers to communicate between processes and allow programs to run concurrently. This communication is achieved using the following intercommunication techniques.

1. Pipes - allows data to flow from one process to another only in one direction typically from the output of one process to the input of another.
2. Although 2 way communication made possible using 2 pipes; 1 in each direction which must have a common process origin. The data from the output process to the input process is buffered until the input process receives it.

Examples :- i) Copying data from a program and pasting it onto another.

ii) Printing a document from the computer system where no data output from the system.

2. Named pipe(s).

In this case the pipe has a specific name.

It can be used by processes that do not have a shared common process origin.

Named pipes are also known as First In First Out (FIFO).

3. Message Queuing.

Message queues allows processes to pass messages between themselves using either a single or several message queues. The system kernel manages each message queue by giving an identity to each message on the queue.

A message queue may be created by one process and then used by multiple processes that read or write messages to the queue.

Application programs (or their processes) create message queues

leaved execution provider major benefit in program structuring.

in system it is possible not only to interleave to overlap them.

and overlapped processes can be viewed as examples process, they both present the same process.

order of execution cannot be predicted it depends on other processes, the way the operating system handles scheduling policies of operating system.

global resources. if 2 processes both make use of it and both perform reads and writes on that in order in which the various reads and writes are done.

the OS to manage the allocation of resources

wait for the OS to simply lock the channel and other processes.

difficult to locate a programming error, because it is not reproducible.

OS have to regulate the order in which data is known as synchronisation.

method of ensuring that when one process is resource (file or variable), the others are excluded.

using co-operating so that only one of the access to a shared resource at a time.

and send and receive messages using API.

4. Semaphores.

They are used to synchronize events between processes. They are integral values greater than or equal to zero.

5. Shared memory.

This allows processes to interchange data through a defined area of memory. example 1 process would write to an area of memory and another could read from it.

To do this writing process must check to see that the read process is reading from the memory at the time and vice versa.

If there are conflicts the other process must wait for the other process to complete.

This is implemented using semaphores where only one process is allowed to access memory at a time.

6. Sockets.

These are typically used to connect over a network between a client and a server.

Although peer to peer connections are also possible.

Sockets are end point of a connection and allow for a standard connection which is computer and OS dependent.

CONCURRENCY.

- A state in which processor exists simultaneously with another process and this is said to be concurrency.

- Principles of concurrency.

In a single multiprocessor multiprogramming system processes are interleaved in time to yield appearance of simultaneous execution even though actual parallel processing is not achieved and certain overhead is involved in switching back and forth between

processors, interleaved execution provides major benefit in processing efficiency and program structuring.

In a multiprocessor system it is possible not only to interleave processes but to overlap them.

Both interleaved and overlapped processes can be viewed as examples of concurrent processes, they both present the same process.

The relative speed of execution cannot be predicted it depends on the activity of other processes, the way the operating system handles interrupts and scheduling policies of operating system.

Following diff

→ sharing of global resources: if 2 processes both make use of a global variable and both perform reads and writes on that variable then the order in which the various reads and writes are executed is critical.

→ It is difficult for the OS to manage the allocation of resources optimally.

- It may be inefficient for the OS to simply lock the channel and prevents its use by other processes.
- It becomes very difficult to locate a programming error, because reports are not usually reproducible.

Synchronization:

The process in which OS have to regulate the order in which processes access data is known as synchronization.

Synchronization issues

→ Mutual exclusion: method of ensuring that when one process is accessing a shared resource (file or variable), the others are excluded from doing the same.

- Two or more processes co-operating so that only one of the processor has access to a shared resource at a time.

2. Race Condition.

- When two concurrent processes are sharing the same storage area.
- A condition in which the behaviour of two or more processes depends on the relative rate at which each process executes its program.

③ Critical Section.

The mechanism for preventing two processes from accessing the same critical area.

A segment of code that cannot be executed while some other process is in a corresponding segment of code.

The fundamental is two or more processes can cooperate by means of simple signal, such that a process can be forced to stop at a specified place until it has received a specified signal.

Any complex coordination requirement can be satisfied by the appropriate structure of signals to transmit a signal via semaphores a process executes a primitive signal.

To receive a signal via semaphores a process executes if the corresponding signal has not yet been transmitted the process is suspended until transmission takes place.

To achieve the desired effect we can view the semaphore as a variable that has an integer value.

① A semaphore may be initialised to a non-negative value.

Wait operation decrements a semaphore value.

If value becomes negative then the process executing the wait is blocked.

The signal operation increments the semaphore value.

Semaphores are operated on by a signal operation which increments the semaphore value and the wait operation. The initial value of a semaphore indicates the number of wait operations that can be performed on the semaphore thus

$$V = I - W + S$$

where I = Initial value of semaphore.

W = number of completed wait operations performed on the semaphore

S = Number of signal operation

V = current value of semaphore which must be greater than or equal to zero

as $V \geq 0$ then $I - W + S \geq 0$ which gives

$$I + S \geq W$$

or

$$W \leq I + S$$

Thus number of weight operations must be less than or equal to the initial value of semaphore plus number of signal operation.

Binary semaphore will have an initial value of 1 ($I=1$) thus:

$$W \leq S + 1.$$